

## Twitter Gender Classification

1. The Problem Statement
2. What we were trying to do/solve/or look for
3. Preparing the Environment
4. The Dataset we chose
5. Reading and Examining the Data
6. Cleaning the Data
7. Pycaret
8. Feature Engineering
9. Models tried
10. Insights
  - a. Comparison to the top-ranking notebook has an accuracy of 48%.
11. Visualizing the Insights
12. Limitations
  - a. PyCaret, Cross-Validation, One-Hot-Encoding Vectorization, String Replacement, etc. all take a lot of computation power and time. Of which, our Google Colab struggles to handle these tasks individually. Having more computation power and time would further allow us to improve our accuracy.

## Twitter Gender Classification

*Machine learning techniques to predict the user's gender on Twitter text data*

This project was produced by MSITM 2021-2022 Team 6:

Shreya Bakshi, Jonathan Cope, Preeti Gupta, Gowtami Khambhampati, Preeti Pinghal, Sneha Reddy, Danqing Wang

### **Introduction:**

The user profile is a persona for a product or service. Gender profiling of unstructured data has several applications in areas such as marketing, advertising, recommendation systems, etc. We can segment the data and understand what drives users, how to attract more users, and how users interact with the service. Now we are going to use the profile information of users on Twitter to predict their gender.

### **The Problem Statement:**

Given this dataset of Twitter conversations, how well can our proposed model predict the gender of the user based on linguistic cues from textual Twitter data.

## Environment:

For this project, we Installed classifiers, necessary packages and imported all the libraries required for visualizations, modeling, automation, vectorization, etc.

We also pip installed pycaret and lgbm.

```
! pip install https://github.com/pandas-profiling/pandas-profiling/archive/master.zip
! pip install pycaret
! pip install lightgbm
```

```
import pandas as pd
import numpy as np
from IPython.display import display
from google.colab import files
!pwd
import os
import re
import regex
import nltk
from nltk.stem import PorterStemmer #Textual data cleaning
nltk.download('stopwords')
from nltk.corpus import stopwords
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams.update({'figure.max_open_warning': 0})
import pylab
from mpl_toolkits.mplot3d import Axes3D # 3D visualization
from matplotlib import pyplot
from mpl_toolkits.axes_grid1 import make_axes_locatable
from scipy import ndimage
from sklearn.feature_extraction import text
from sklearn import model_selection
from sklearn.naive_bayes import MultinomialNB # Naive Bayes model
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

import seaborn as sns
from pandas_profiling import ProfileReport # This library is great for seeing what is missing and what needs to be cleaned up
from bokeh.plotting import output_notebook, figure, show #graph
from bokeh.layouts import gridplot
from bokeh.models import ColumnDataSource

from collections import Counter
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from lightgbm import LGBMClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import ComplementNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from sklearn.pipeline import Pipeline
from sklearn.neighbors import KNeighborsClassifier
```

## Data Source:

The data has been extracted from Kaggle from the link below.

<https://www.kaggle.com/crowdflower/twitter-user-gender-classification>

Read the CSV file:


```
[3] twitter_data = files.upload()
data = pd.read_csv("gender-classifier-DFE-791531.csv", encoding='latin-1')
```

[Browse...](#) gender-classifier-DFE-791531.csv

gender-classifier-DFE-791531.csv(text/csv) - 8176739 bytes, last modified: n/a - 100% done

Saving gender-classifier-DFE-791531.csv to gender-classifier-DFE-791531 (1).csv

## Shape of the dataset:

 data.shape

(20050, 26)

Dataset consists of 20050 rows and 26 columns. Out of 26 columns, we use 25 predictor variables and 1 target variable which is 'gender'.

## Columns and their associated data types:

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20050 entries, 0 to 20049
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   _unit_id                             20050 non-null  int64
1   _golden                              20050 non-null  bool
2   _unit_state                           20050 non-null  object
3   _trusted_judgments                   20050 non-null  int64
4   _last_judgment_at                   20000 non-null  object
5   gender                               19953 non-null  object
6   gender:confidence                    20024 non-null  float64
7   profile_yn                           20050 non-null  object
8   profile_yn:confidence                20050 non-null  float64
9   created                              20050 non-null  object
10  description                           16306 non-null  object
11  fav_number                           20050 non-null  int64
12  gender_gold                           50 non-null     object
13  link_color                           20050 non-null  object
14  name                                  20050 non-null  object
15  profile_yn_gold                       50 non-null     object
16  profileimage                          20050 non-null  object
17  retweet_count                         20050 non-null  int64
18  sidebar_color                         20050 non-null  object
19  text                                  20050 non-null  object
20  tweet_coord                           159 non-null    object
21  tweet_count                           20050 non-null  int64
22  tweet_created                         20050 non-null  object
23  tweet_id                              20050 non-null  float64
24  tweet_location                       12566 non-null  object
25  user_timezone                         12252 non-null  object
dtypes: bool(1), float64(3), int64(5), object(17)
memory usage: 3.8+ MB
```

## Analyzing Data source:

### Features of dataset:

```
[ ] data.keys()
```

```
Index(['_unit_id', '_golden', '_unit_state', '_trusted_judgments',
      '_last_judgment_at', 'gender', 'gender:confidence', 'profile_yn',
      'profile_yn:confidence', 'created', 'description', 'fav_number',
      'gender_gold', 'link_color', 'name', 'profile_yn_gold', 'profileimage',
      'retweet_count', 'sidebar_color', 'text', 'tweet_coord', 'tweet_count',
      'tweet_created', 'tweet_id', 'tweet_location', 'user_timezone'],
      dtype='object')
```

## Exploring Target variable - 'Gender':

```
data['gender'].value_counts()
```

```
female    6700
male      6194
brand      5942
unknown   1117
Name: gender, dtype: int64
```

## Null values in the dataset:

```
data.isnull().sum()
```

```
_unit_id          0
_golden           0
_unit_state       0
_trusted_judgments 0
_last_judgment_at 50
gender            97
gender:confidence 26
profile_yn        0
profile_yn:confidence 0
created           0
description       3744
fav_number        0
gender_gold       20000
link_color        0
name              0
profile_yn_gold   20000
profileimage      0
retweet_count     0
sidebar_color     0
text              0
tweet_coord       19891
tweet_count       0
tweet_created     0
tweet_id          0
tweet_location    7484
user_timezone     7798
dtype: int64
```

## Data Cleaning:

Data cleansing is a very crucial step in the overall data preparation process and it is the process of analyzing, identifying, and correcting messy, raw data.

We started our data cleaning by dropping unnecessary features

#### Drop unnecessary columns/features

```
data.drop (columns = ['_unit_id',
                      '_last_judgment_at',
                      'user_timezone',
                      'tweet_coord',
                      'tweet_created',
                      'tweet_id',
                      'tweet_location',
                      'profileimage',
                      'created',
                      'name'], inplace = True)

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20050 entries, 0 to 20049
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   _golden                20050 non-null  bool
1   _unit_state            20050 non-null  object
2   _trusted_judgments     20050 non-null  int64
3   gender                 19953 non-null  object
4   gender:confidence      20024 non-null  float64
5   profile_yn             20050 non-null  object
6   profile_yn:confidence  20050 non-null  float64
7   description            16306 non-null  object
8   fav_number             20050 non-null  int64
9   gender_gold            50 non-null     object
10  link_color              20050 non-null  object
11  profile_yn_gold         50 non-null     object
12  retweet_count           20050 non-null  int64
13  sidebar_color           20050 non-null  object
14  text                    20050 non-null  object
15  tweet_count             20050 non-null  int64
dtypes: bool(1), float64(2), int64(4), object(9)
memory usage: 2.3+ MB
```

#### Getting rid of gender type - “unknown”:

```
unknown_items_idx = data[data['gender'] == 'unknown'].index
data.drop (index = unknown_items_idx, inplace = True)
data['gender'].value_counts()
```

```
female    6700
male       6194
brand       5942
Name: gender, dtype: int64
```

#### Getting rid of rows where column “profile\_yn” is no:

```
[20] drop_items_idx = data[data['profile_yn'] == 'no'].index
data.drop (index = drop_items_idx, inplace = True)
data.drop (columns = ['profile_yn', 'profile_yn:confidence', 'profile_yn_gold'], inplace = True)
```

#### Removing Stop words and cleaning the Text:

Here we are using functions `preprocessor()`, `remove_dup_whitespace()`, `tokenizer_porter()`, `clean_tweet`, `has_nan` to clean, stem and tokenize the text

```

stop = stopwords.words('english')
porter = PorterStemmer()

def preprocessor(text):
    #Return a cleaned version of text, but keeping the emoticons
    text = re.sub('<[>]*>', '', text) # Remove HTML markup
    text = re.sub('http.*', ' ', text) # Remove url tokens
    return text

def remove_dup_whitespace(text):
    #This function removes duplicated whitespaces of a string
    return re.sub('\s{2,}', ' ', text) #return text

def tokenizer_porter(text):
    # This function tokenize and also perform stemming
    return [porter.stem(word) for word in text.lower().split()]

def clean_tweet(text):
    #This function tokenizes whole tweet into tokens, clean it, remove stopwords and combine back as a tweet, this function com
    clean = ""
    tokens = tokenizer_porter(text)
    for token in tokens:
        if len(token) > 1:
            if token not in stop:
                clean += preprocessor(token) + " "
    #print(data)
    return clean
    return remove_dup_whitespace(clean)

def has_nan(X):
    """
    Input: Dataframe
    This func check if the features of a Dataframe has missing values or not
    """
    X_ = X.isnull()
    X_ = X_.add_suffix('_has_nan')
    return X_

has_nan_df = has_nan(data[['description']])
data = pd.concat([data, has_nan_df], axis=1)

data['description'].fillna("", inplace=True) # Fill NaN with empty string

/usr/local/lib/python3.7/dist-packages/pandas/core/series.py:4536: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
downcast=downcast,

```

```
data[['text', 'description']] = data[['text', 'description']].applymap(clean_tweet)
```

## Changing the text to lower characters:

```

def clean(review1):
    descrip = re.sub('[a-zA-Z]', ' ', review1)
    review1 = review1.lower()
    return review1

data['text_Cleaned'] = pd.DataFrame(data['text'].apply(lambda y: clean(y)))
data.head(5)

```

	_golden	_unit_state	_trusted_judgments	gender	gender:confidence	description	fav_number	gender_gold	link_color	retweet_count	sidebar_color	text	tweet_count	description_has_nan	text_Cleaned
0	False	finalized	3	male	1.0000	sing rhythm.	0	NaN	08C2C2	0	FFFFFF	robbi respond critic win eddi edward #worldtit...	110964	False	robbi respond critic win eddi edward worldtit...
1	False	finalized	3	male	1.0000	i'm author novel fill famili drama romance.	68	NaN	0084B4	0	CODEED	Ülit felt like friend wa live stori themü ...	7471	False	Ülit felt like friend wa live stori themü ...
2	False	finalized	3	male	0.6625	loui whine squeal	7696	NaN	AB88C2	1	CODEED	absolut ador loui start song hit hard feel good	5617	False	absolut ador loui start song hit hard feel good
3	False	finalized	3	male	1.0000	mobli guy. 49ers, shazam, google, kleiner perk...	202	NaN	0084B4	0	CODEED	hi @jordanspieth look url use @fittt?i typic s...	1693	False	hi @jordanspieth look url use @fittt?i typic s...
4	False	finalized	3	female	1.0000	ricki wilson best frontmankais chief best ban...	37318	NaN	3B94D9	0	0	watch neighbour sky+ catch neighbour! xxx _lo4_...	31462	False	watch neighbour sky+ catch neighbour! xxx _lo4_...

```

data['text_Cleaned'] = data['text'].str.replace('[A-Za-z0-9 ]+', '')
data['text_Cleaned'] = data['text_Cleaned'].str.lower()
data.head()

```

	gender	gender:confidence	description	fav_number	link_color	retweet_count	sidebar_color	text	tweet_count	description_has_nan	text_Cleaned
0	male	1.0000	sing rhythm.	0	08C2C2	0	FFFFFF	robbi respond critic win eddi edward #worldtit...	110964	False	robbi respond critic win eddi edward worldtit...
1	male	1.0000	i'm author novel fill famili drama romance.	68	0084B4	0	CODEED	Ülit felt like friend wa live stori themü ...	7471	False	it felt like friend wa live stori them retr...
2	male	0.6625	loui whine squeal	7696	AB88C2	1	CODEED	absolut ador loui start song hit hard feel good	5617	False	absolut ador loui start song hit hard feel good
3	male	1.0000	mobli guy. 49ers, shazam, google, kleiner perk...	202	0084B4	0	CODEED	hi @jordanspieth look url use @fittt?i typic s...	1693	False	hi jordanspieth look url use ifttt? typic see a...
4	female	1.0000	ricki wilson best frontmankais chief best ban...	37318	3B94D9	0	0	watch neighbour sky+ catch neighbour! xxx _lo4_...	31462	False	watch neighbour sky catch neighb xxx xxx

We have removed the column description\_has\_nan as it has no significance in predicting the gender which is our goal here.

```
data.drop(['text', 'description_has_nan'],axis=1,inplace=True)
data.head()
```

	gender	gender:confidence	description	fav_number	link_color	retweet_count	sidebar_color	tweet_count	text_Cleaned
0	male	1.0000	sing rhythm.	0	08C2C2	0	FFFFFF	110964	robbi respond critic win eddi edward worldtit...
1	male	1.0000	i'm author novel fill famli drama romance.	68	0084B4	0	C0DEED	7471	it felt like friend wa live stori them retir...
2	male	0.6625	loui whine squeal	7696	ABB8C2	1	C0DEED	5617	absolut ador loui start song hit hard feel good
3	male	1.0000	mobil guy, 49ers, shazam, google, kleiner perk...	202	0084B4	0	C0DEED	1693	hi jordanspleth look uri use iftt typic see a...
4	female	1.0000	ricki wilson best frontman/kais chief best ban...	37318	3B94D9	0	0	31462	watch neighbour sky catch neighb xxx xxx

## PyCaret Machine Learning

PyCaret is an automated tool in Python that allows (with time) to create insights on what kind of data is being collected (categorical, numerical, boolean, etc.), and give the user a chance to make adjustments and tweaks to the auto-generated assumptions. The tool will then compile a variety of models using default hyperparameters to provide results such as accuracy and AUC. The tool can be further used to boost/ensemble/predict.

On the cleaned data without much feature engineering, we were able to find that out of models (Light Gradient Boosting, SVM, Random Forest, Logistic Regression, Ridge, Naive Bayes, Linear Discriminant Analysis, Extra Trees, Gradient Boosting, Quadratic Discriminant Analysis, Ada Boost, Decision Tree, K Neighbors, Dummy) and it tells us here the accuracy was 58%.

```
compare_models(sort = 'Accuracy', fold = 5)
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>lightgbm</b>	Light Gradient Boosting Machine	0.5808	0.7675	0.5788	0.5929	0.5816	0.3674	0.3710	1.956
<b>svm</b>	SVM - Linear Kernel	0.5701	0.0000	0.5685	0.5769	0.5691	0.3517	0.3548	3.112
<b>rf</b>	Random Forest Classifier	0.5675	0.7510	0.5680	0.5673	0.5665	0.3495	0.3502	15.542
<b>lr</b>	Logistic Regression	0.5661	0.7513	0.5655	0.5714	0.5669	0.3465	0.3477	9.128
<b>ridge</b>	Ridge Classifier	0.5659	0.0000	0.5647	0.5756	0.5675	0.3457	0.3477	0.622
<b>nb</b>	Naive Bayes	0.5655	0.7326	0.5628	0.5765	0.5647	0.3440	0.3481	0.256
<b>lda</b>	Linear Discriminant Analysis	0.5625	0.7463	0.5610	0.5814	0.5663	0.3401	0.3429	5.518
<b>et</b>	Extra Trees Classifier	0.5612	0.7424	0.5620	0.5593	0.5594	0.3405	0.3411	20.846
<b>gbc</b>	Gradient Boosting Classifier	0.5481	0.7451	0.5392	0.6063	0.5360	0.3119	0.3409	37.750
<b>qda</b>	Quadratic Discriminant Analysis	0.5295	0.6996	0.5257	0.5821	0.5201	0.2887	0.3117	5.498
<b>ada</b>	Ada Boost Classifier	0.5160	0.6957	0.5041	0.5993	0.4667	0.2596	0.3219	3.946
<b>dt</b>	Decision Tree Classifier	0.5118	0.6424	0.5125	0.5097	0.5101	0.2661	0.2665	4.990
<b>knn</b>	K Neighbors Classifier	0.4630	0.6423	0.4599	0.4796	0.4635	0.1888	0.1917	78.654
<b>dummy</b>	Dummy Classifier	0.3562	0.5000	0.3333	0.1269	0.1871	0.0000	0.0000	0.064

```
LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
importance_type='split', learning_rate=0.1, max_depth=-1,
min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,
n_estimators=100, n_jobs=-1, num_leaves=31, objective=None,
random_state=101, reg_alpha=0.0, reg_lambda=0.0, silent='warn',
subsample=1.0, subsample_for_bin=200000, subsample_freq=0)
```

## Feature Engineering:

To get the most out of our PyCaret Auto Modeling, we had to perform a few tasks of Feature Engineering. For example, we began by reindexing the data (since it would skip numbers, which would affect adding back the Y label after one-hot-encoding). Then, we pursued one hot encoding, starting with 500 columns for the “description” column and 500 for the “text\_Cleaned” column. We would rename the second set of columns “500”-“999” so that they wouldn't clash with the 0-499 set.

```
data_copy = data.copy().reset_index().drop('index', axis=1)

cv = CountVectorizer(max_features = 500)
pyc = cv.fit_transform(data_copy['description']).toarray()
pyc1=cv.fit_transform(data_copy['text_Cleaned']).toarray()
```

```
P1=pd.DataFrame(pyc)
P2=pd.DataFrame(pyc1)
# We have to rename the second set of columns because two sets of the same names will crash
P2.columns = [x for x in range(500, 1000)]
```

Since PyCaret would not mix the string ‘brand’ in well, with the gender integers (0, 1), we translated ‘brand’ as a ‘2.’ We also converted the datatypes for the table as “int8” so that the computation would run faster than the default int64.



```
[ ] # we shouldn't mix ints and string labels together so lets fix that
    for gen in data_copy['gender']:
        if gen == 'brand':
            data_copy['gender'].replace({'brand': '2'}, inplace=True)
```

```
data_copy['gender']
```

```
0      1
1      1
2      1
3      1
4      0
..
18831   0
18832   1
18833   1
18834   0
18835   0
Name: gender, Length: 18836, dtype: object
```

```
[ ] train_for_pycaret=pd.concat([P2,P1],join='outer',axis=1)
    train_for_pycaret['gender'] = data_copy['gender']
    train_for_pycaret=train_for_pycaret.fillna(0)
    train_for_pycaret = train_for_pycaret.astype('int8')
    train_for_pycaret.columns = [str(x) for x in train_for_pycaret.columns]
    train_for_pycaret
```

We then noticed that if we improve our one-hot-encoding vectorizer from 500 to 1500, the accuracy improves from 56 to 63.14% with Extra Trees Classifier

```
cv = CountVectorizer(max_features = 1500)
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>et</b>	Extra Trees Classifier	0.6314	0.7973	0.6371	0.6286	0.6282	0.4405	0.4420	33.950
<b>rf</b>	Random Forest Classifier	0.6304	0.7977	0.6325	0.6328	0.6275	0.4361	0.4396	20.202
<b>lr</b>	Logistic Regression	0.6303	0.7949	0.6339	0.6343	0.6310	0.4372	0.4381	21.052
<b>lightgbm</b>	Light Gradient Boosting Machine	0.6247	0.7998	0.6247	0.6333	0.6243	0.4262	0.4297	2.486

After this, we improved our accuracy score further by including the Link Color (categorical) data, along with the Favorite Number (numerical) data.

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>lightgbm</b>	Light Gradient Boosting Machine	0.6562	0.8289	0.6617	0.6558	0.6553	0.4782	0.4788	2.640
<b>rf</b>	Random Forest Classifier	0.6521	0.8225	0.6537	0.6555	0.6497	0.4693	0.4729	18.362
<b>et</b>	Extra Trees Classifier	0.6454	0.8129	0.6480	0.6479	0.6436	0.4596	0.4622	32.608
<b>gbc</b>	Gradient Boosting Classifier	0.6273	0.8068	0.6291	0.6268	0.6243	0.4327	0.4349	94.150

Finally, we decided to go a step further with our Link Color data (which was previously in HEX format (ex. #AFFF00)). We converted this data into RGB Values (Red, Green, Blue values of 0-255 inclusive). Once we had these values in a usable format, we converted them into Color names using PIL's ImageColor library. If the name included "light" or "dark" substrings, we would remove that portion of the name. That way we would have multiple values using the standard labels "red," "blue," "gold," etc. regardless of whether they were light reds, dark blues, etc. This also made the data more categorical, and we also included the favorite color column as categorical data. **This brought our accuracy score to 65.99%.**

Based off of fiatjaf's comment and modified: <https://stackoverflow.com/questions/9694165/convert-rgb-color-to-english-color-name-like-green-with-python>

```
[ ] import matplotlib.colors as mc
    mycss4list = mc.CSS4_COLORS

    # For HEX to RGB
    from PIL import ImageColor

    def getColorName(hex_input):
        min_colors = {}
        for name, hex in mycss4list.items():
            r, g, b = ImageColor.getcolor(hex, "RGB")
            r_input, g_input, b_input = ImageColor.getcolor(hex_input, "RGB")
            rd = (r - r_input) ** 2
            gd = (g - g_input) ** 2
            bd = (b - b_input) ** 2
            min_colors[(rd + gd + bd)] = name
        #print(min_colors)
        colorName = min_colors[min(min_colors.keys())]
        colorName = colorName.replace("dark", "")
        colorName = colorName.replace("light", "")
        return colorName

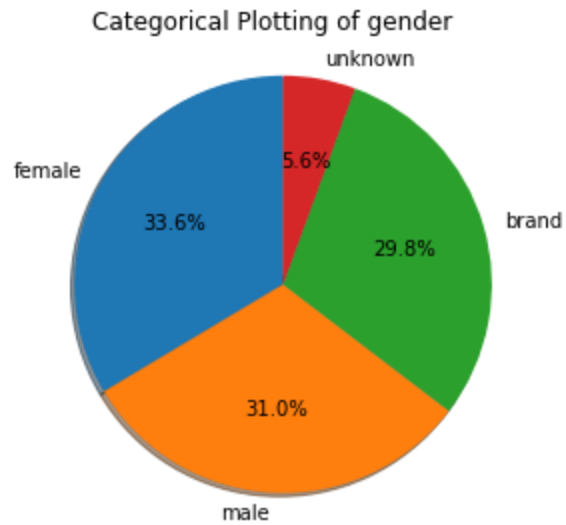
[ ] print(getColorName("#990003"))

red
```

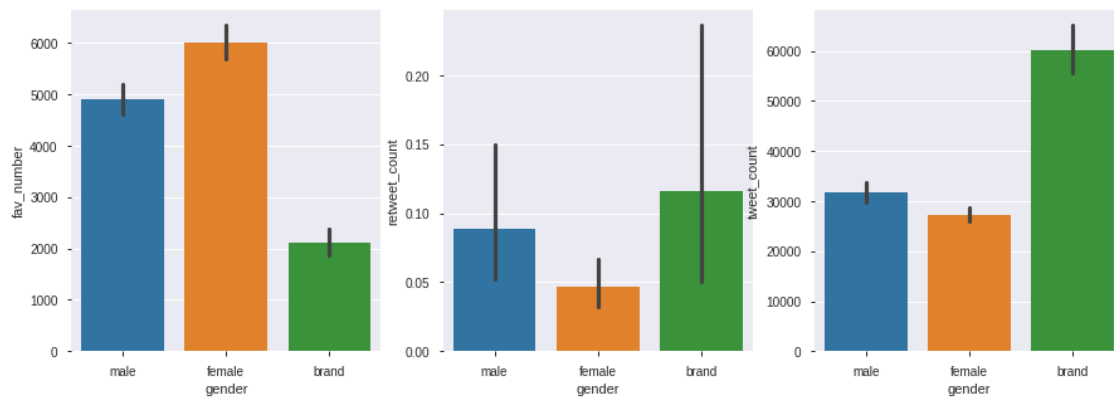
	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
lr	Logistic Regression	0.6599	0.8211	0.6638	0.6636	0.6610	0.4829	0.4834	49.698
svm	SVM - Linear Kernel	0.6369	0.0000	0.6413	0.6420	0.6386	0.4484	0.4490	10.040

## Visualizing the Data:

### Gender distribution

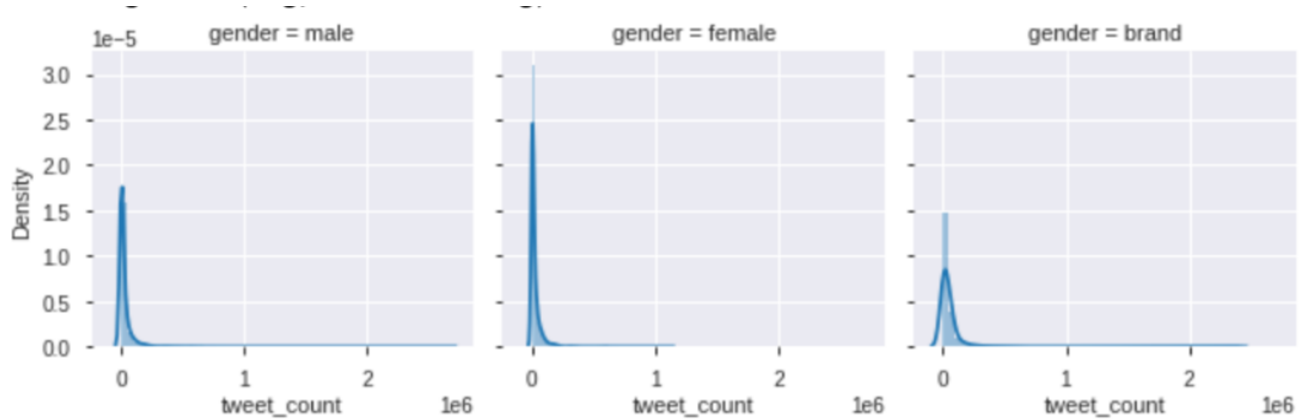


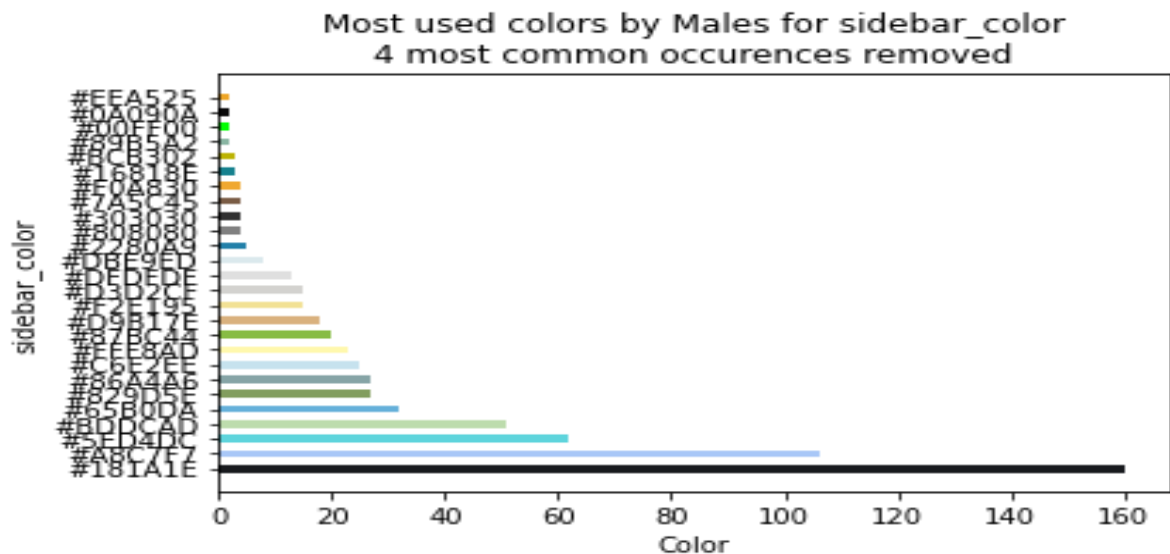
Subplots for fav\_number, retweet\_count, tweet\_count for all 'gender' types:



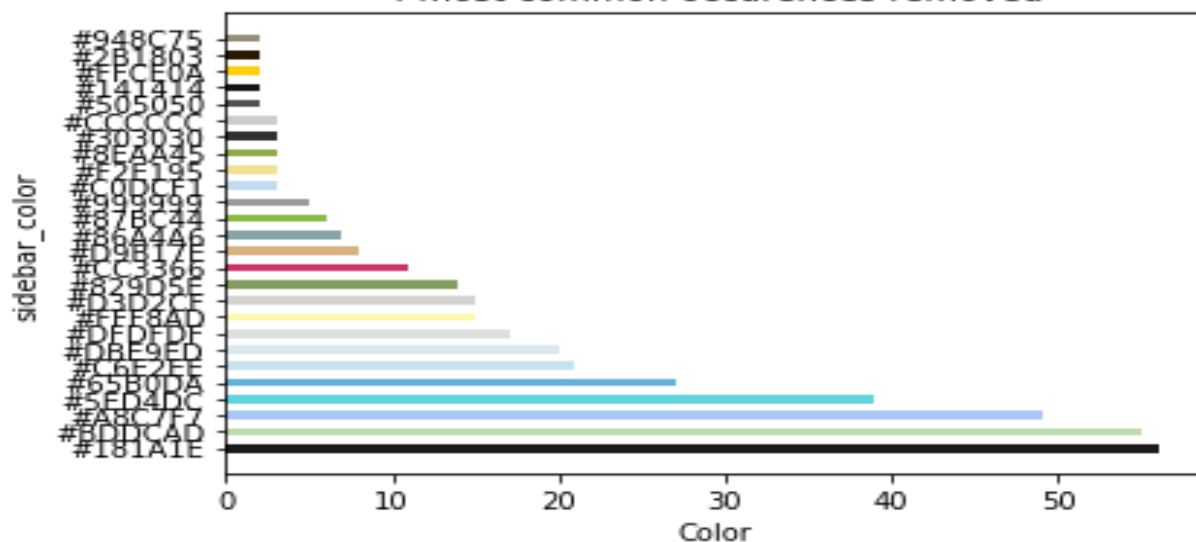
From the figure above we notice that the retweet\_count and tweet\_count for brands are higher when compared to others.

Density Graph of Tweet count vs Gender

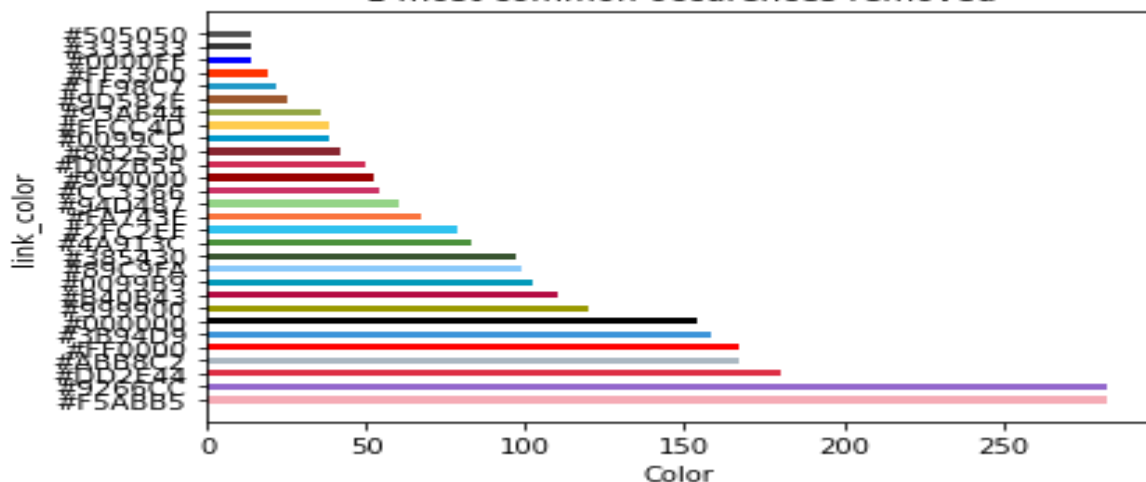




Most used colors by Brands for sidebar\_color  
4 most common occurrences removed

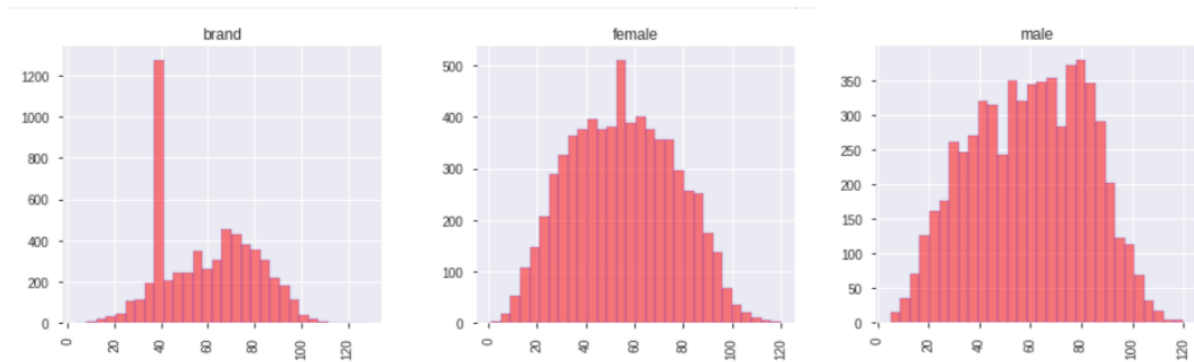


Most used colors by Females for link\_color  
1 most common occurrences removed





## Text length vs Gender



The above image shows a comparison of text lengths and how they differ among the gender and brand. We can see data follows a normal distribution trend here.

**Word cloud of “text” column based on gender:**

### Word Cloud of Gender: Male



### Word Cloud of Gender: Female



### Compare and contrast two classics

### Label Encoding:



```

from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()
y = encoder.fit_transform(data['gender'])

# split the dataset in train and test
X = data['text_Cleaned']
# Stratify will create a train set with the same class balance than the original set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0, stratify=y)

```

## **Splitting train and test data:**

The train and test data is split into 70% and 30%

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0, stratify=y)

```

## **Training data on different ML models:**

We have tried to implement the different prediction models just by using text and non-text and text columns without involving any sentiment analysis with the textual data.

Features used for the predictions:

**Link Color:** It has non-text data and indicates the link color on the profile

**Description:** The user's profile description

**Text:** Text of a random one of the user's tweets

We found these attributes to provide useful information regarding gender classification.

**Label Encoder:** Before the prediction, we have encoded the target column to 0-1

## **Predictions using non-text column- Link Color**

Twitter allows customizing and personalizing the account by changing the colors of the links or the sidebars, and we expect people from different genders to have different behaviors in how they personalize their page.

Hence, we have used link color as the feature for this prediction.

```

def model_test(model,X_train,y_train,X_test,y_test, full_voc, displayResults = True, displayColors = False, featureIntent = 'text'):
    switcher = {

        'link_color' : "theme color",
    }
    featureText = switcher.get(featureIntent, '')

    model.fit(X_train,y_train)
    y_pred = model.predict(X_test)

    # compute MSE
    mse = metrics.mean_squared_error(y_test,y_pred)
    print('mse: {:.4f}'.format(mse))

    # Prints the accuracy of the gender prediction
    acc = model.score(X_test,y_test)
    print('score: ', acc)

    import matplotlib.pyplot as plt
    import sklearn
    conf = sklearn.metrics.confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(5,5))
    sns.heatmap(conf, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues');
    plt.ylabel('Actual label');
    plt.xlabel('Predicted label');
    plt.imshow(conf, cmap='binary', interpolation=None)
    plt.show()

    return model, acc

```

```

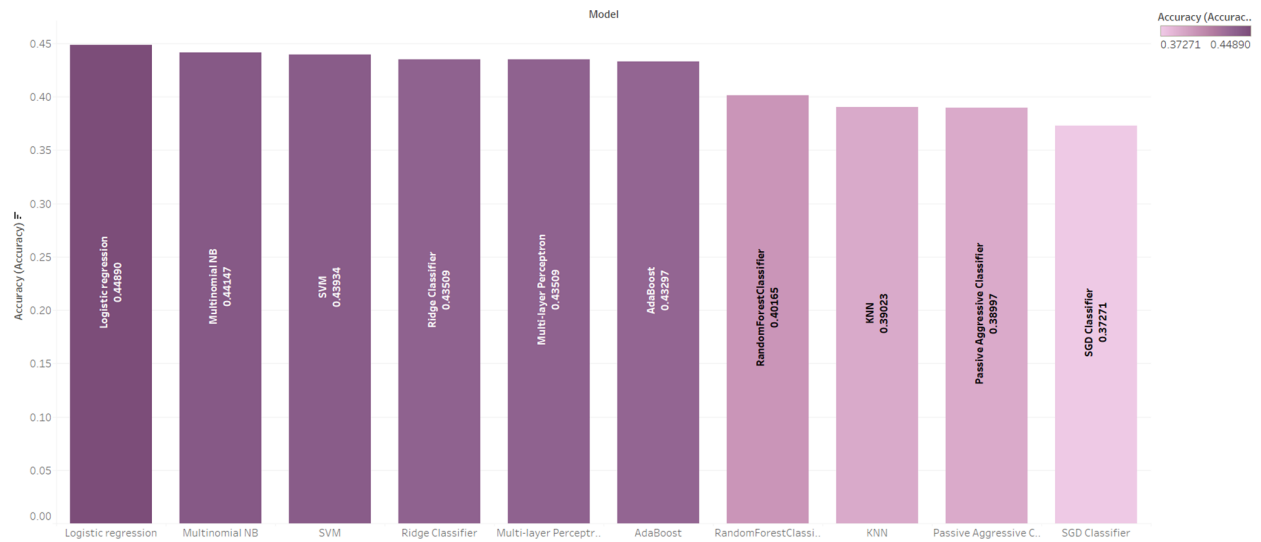
def compute_bag_of_words(text):
    vectorizer = CountVectorizer()
    vectors = vectorizer.fit_transform(text)
    vocabulary = vectorizer.get_feature_names()
    return vectors, vocabulary

def predictors(df, feature, model, modelname, displayResults = True, displayColors = False):
    print('Testing', modelname, 'model for gender prediction using', feature)
    full_bow, full_voc = compute_bag_of_words(df[feature])
    X = full_bow
    y = LabelEncoder().fit_transform(df['gender'])
    # Create Training and testing sets.
    n,d = X.shape
    test_size = n // 5
    print('Split: {} testing and {} training samples'.format(test_size, y.size - test_size))
    perm = np.random.permutation(y.size)
    X_test = X[perm[:test_size]]
    X_train = X[perm[test_size:]]
    y_test = y[perm[:test_size]]
    y_train = y[perm[test_size:]]
    print('model: ', modelname)
    model, acc = model_test(model,X_train,y_train,X_test,y_test, full_voc, displayResults = displayResults, displayColors = displayColors, featureIntent = feature)
    return model, full_voc, acc

```

## Accuracy Scores

## Non-Text Vs Accuracy



Sum of Accuracy (Accuracy) for each Model. Color shows sum of Accuracy (Accuracy). The marks are labeled by Model and sum of Accuracy (Accuracy). The view is filtered on Model, which has multiple members selected.

Logistic Regression gave the highest accuracy followed by Multinomial NB SGD has the lowest scores, whereas Adaboost also gave a low accuracy score.

## Predictions using text column- Text

We have cleaned the text column before doing the prediction in order to get rid of noisy data.

We have used TfidfVectorizer to calculate the TF-IDF values and understand the importance and weightage of a word in the text.

```

def classification_modeling(X_train, X_test, y_train, y_test, text_feature=False):
    """
    This function iterates different possible models
    and return corresponding accuracy

    Args:
        text_feature: Whether the model handles text features or not

    Return: The best fitted model
    """
    clf_dict = {'lr': linear_model.LogisticRegression(multi_class='ovr', random_state=0),
                'rf': ensemble.RandomForestClassifier(n_estimators = 50, random_state=0),
                'svm': SVC(kernel = 'rbf', probability=True),
                'nb': ComplementNB(),
                'ridge': linear_model.RidgeClassifier(),
                'sgd': linear_model.SGDClassifier(),
                'passaggre': linear_model.PassiveAggressiveClassifier(),
                'NB': naive_bayes.MultinomialNB(),
                'NN': neural_network.MLPClassifier(),
                'Knn': neighbors.KNeighborsClassifier(n_neighbors=5, weights='distance', algorithm='auto'),
                'Adaboost': ensemble.AdaBoostClassifier()
                }
    result_dict = dict.fromkeys(clf_dict, None)
    pred_dict = dict.fromkeys(clf_dict, None)

    modelNamesList = [
        'LogisticRegression',
        'RandomForestClassifier',
        'SVM',
        'ComplementNB',
        'RidgeClassifier',
        'SGDClassifier',
        'PassiveAggressiveClassifier',
        'MultinomialNB',
        'MLPClassifier',
        'KNN',
        'Adaboost'
    ]

    acc_color = np.zeros(len(modelNamesList))

    acc_val=[]
    for clf_key in clf_dict:
        if text_feature == True:
            tfidf = TfidfVectorizer()
            clf = Pipeline([('vect', tfidf),
                            ('clf', clf_dict[clf_key])])
        else:
            clf = clf_dict[clf_key]
            clf.fit(X_train, y_train)
            predictions = clf.predict(X_test)

```

```

acc = accuracy_score(y_test,predictions)
result_dict[clf_key] = acc
pred_dict[clf_key] = predictions
print('Fitting ' + clf_key + ' - Acc:', acc)
acc_val.append(acc);

import matplotlib.pyplot as plt
import sklearn
conf = sklearn.metrics.confusion_matrix(y_test,predictions)
plt.figure(figsize=(5,5))
sns.heatmap(conf, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'icefire');
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.imshow(conf, cmap='binary', interpolation='None')
plt.show()
#print('Confusion matrix:\n',confusion_matrix(y_test,predictions))
#print('-'*40)
#print(acc_val)

win_clf = max(result_dict, key=lambda key: result_dict[key])
print("Win classifier: ", win_clf, "- Acc: ",result_dict[win_clf])
fig, ax1 = plt.subplots(figsize=(8,8))
ax1.set_xlim([0, 1])
#bar_width = 0.5
#plt.figure(figsize=(8,8))
model_number = np.arange(len(modelNamesList))+1
rects1 = plt.barh(model_number, acc_val,color = '#D95319')
plt.yticks(model_number,modelNamesList)
plt.xlabel('Accuracy with text')
plt.ylabel('Model')
plt.title('Accuracy of the different Classifiers')
plt.tight_layout()
plt.show()
return np.asarray(pred_dict[win_clf])

```

## Accuracies:

Text Vs Accuracy



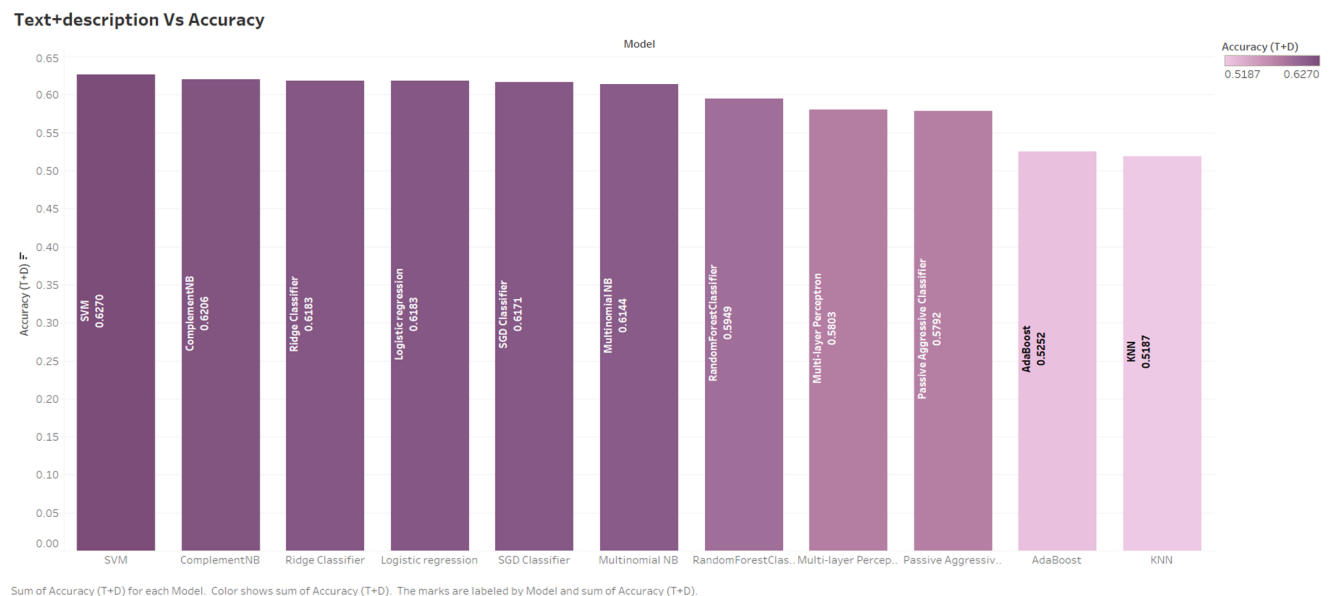
Support vector machines (SVM) have the highest accuracy score followed by SGD- stochastic gradient descent.

Overall accuracy score increased compared to scores from using just link color feature but still, it is between 47%-53% and has not improved significantly.

## Predictions using text + description column

```
| data['text_description'] = data['text_Cleaned'].str.cat(data['description'], sep=' ')\n\n| X = data['text_description']\n  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0, stratify=y)\n  #In the code line above, stratify will create a train set with the same class balance than the original set\n\n  X_train.head()\n  X_train.isnull().values.any() # Check if any null values, True if there is at least one.\n\n  False\n\n| best_text_preds = classification_modeling(X_train, X_test, y_train, y_test, text_feature=True)
```

We have concatenated the text and the description column and then have used the same code as described above



SVM again has the highest score followed by Complement Naive Bayes and the accuracy scores have increased significantly when we are combining the text and the description column. We have better prediction chances using the user's profile description and the text they are tweeting. There is no single prediction model which performs well in all the cases; however, we see that overall SVM has higher accuracy predicting the gender compared to other models.

Note- Prediction graphs generated using tableau , based on the data from the prediction models

By using boosting classifiers like LGBM Classifier, We are getting an accuracy of 56%

```

lgbmodel = LGBMClassifier(max_depth=5)
lgbmodel.fit(X_train, y_train)
y_pred_lgbm = lgbmodel.predict(X_test)
accuracy_lgbm = accuracy_score(y_test, y_pred_lgbm)
print("Accuracy: %.2f%%" %(accuracy_lgbm * 100.0))

```

Accuracy:56.71%

By using SVM Classifier, We are getting an accuracy of 62.7%

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
lr	Logistic Regression	0.6599	0.8211	0.6638	0.6636	0.6610	0.4829	0.4834	49.698

With PyCaret, and feature engineering (explained previously), our highest possible accuracy given our computation limits was 65.99%.